

Feasibility (and what will *actually* make this publishable)

What you described is **feasible** in IsaacLab/Isaac Sim: you can train an RL pick-and-place policy using **RGB-D + proprio**, then evaluate it under **vision modality dropout**, then train a second policy that additionally observes **gripper force / contact signals**, and compare robustness.

However, there are two “gotchas” that will determine whether your paper reads like a solid contribution versus an obvious demo:

1. **Force sensing will not help if vision dropout happens “too early.”**
If the camera drops out **before the robot has localized the object**, force doesn’t provide object position—so both models will fail similarly.
→ To make the force sensor matter, your evaluation should include **phase-aware dropout** (e.g., during grasp closure / after grasp) and/or a task where the **goal is known but the object must be acquired**.
2. **If you don’t train with any modality corruption, the force-sensor model may still ignore force.**
Even if force is present in the observation, the easiest signal is usually vision. If vision is always perfect during training, the network can learn a “vision-only” solution and treat force as noise.
→ For a publication-quality robustness claim, you should include **at least one robust-training baseline** (vision dropout/noise during training), otherwise reviewers will say “you could have just trained with dropout.”

Recommended change (be specific)

Keep your original two-model story, but **add one extra baseline** so your conclusion is defensible:

- **M1 (baseline):** RGB-D + proprio, trained with clean vision
- **M2 (your proposed):** RGB-D + proprio + force, trained with clean vision
- **M3 (recommended baseline):** RGB-D + proprio, trained with *vision corruption/dropout*
- **M4 (best-method, optional but strong):** RGB-D + proprio + force, trained with *vision corruption/dropout*

This lets you separate:

- “force helps” vs.
- “dropout training helps” vs.
- “force + dropout training is best.”

If you only do M1 vs M2, you risk a weak or inconsistent effect.

The paper's core claim (make it crisp)

A clean publication framing:

Claim: Adding gripper contact/force sensing improves robustness of vision-based RL manipulation under intermittent vision failures; robustness depends strongly on *when* and *for how long* vision is missing.

Your contribution becomes:

1. A reproducible IsaacLab benchmark: RGB-D pick-and-place with controlled vision failures
 2. A multi-modal policy with force/contact input
 3. A systematic study of **dropout duration + dropout timing** sensitivity
-

Experimental design that won't collapse under review

Task: simple but not trivial

Use a single object pick-and-place with randomized object pose:

- Robot: **Franka Panda** (common, good controllers, easy)
 - Object: cube or cylinder (start with cube)
 - Table scene: uncluttered (for “simple”)
 - Initial object pose: random in a rectangle
 - Goal: a bin/target region at a known pose
- Important:** provide goal pose as a **command input** (low-dimensional), not only visually. This makes “vision dropout after grasp” meaningful.

Episode structure

- Control frequency: **20 Hz**
- Episode length: **250 steps** (12.5 s)
- Reset randomization each episode

Success definition (write it clearly)

Success at end of episode or terminal condition if all true:

- Object position within goal radius:

$$\|p_{obj} - p_{goal}\| < 3 \text{ cm}$$

- Object height near goal surface (not hovering): e.g. $|z_{\text{obj}} - z_{\text{goal}}| < 2 \text{ cm}$ $|z_{\text{obj}} - z_{\text{goal}}| < 2 \text{ cm}$
 - Object linear speed small (stable): $\|\mathbf{v}_{\text{obj}}\| < 0.05 \text{ m/s}$ $\|\mathbf{v}_{\text{obj}}\| < 0.05 \text{ m/s}$
 - Gripper open (released), optional but recommended
-

Observations (what each model sees)

Common proprio + command (both models)

Include:

- Joint positions \mathbf{q}_{qq} (7)
- Joint velocities $\dot{\mathbf{q}}$ (7)
- Gripper width / finger positions (1–2)
- End-effector pose (position + orientation) OR (position + quaternion) (7)
(This is still “proprio-derived,” not external perception.)
- Previous action (helps stability) (action_dim)
- **Goal pose command** in robot base frame: $\mathbf{p}_{\text{goal}} \mathbf{p}_{\text{goal}}$ (3) (and optionally yaw)

This matters because with vision dropout, the robot still needs to know where “place” is.

Vision modality (RGB-D)

- One camera (choose one and justify):
 - **Wrist-mounted** camera: best for grasp; more sensitive to occlusion
 - **Static overhead**: best for global localization; less realistic for many setups
- Resolution: **84×84** or **96×96** (start small; render cost dominates)
- Inputs:
 - RGB normalized to [0,1]
 - Depth clipped to [0, d_{max}] and normalized (e.g., $d_{\text{max}}=1.5$) $= 1.5 \text{ d}_{\text{max}} = 1.5 \text{ m}$)
- Optional: 2-frame stack (gives motion cues)

Force / contact modality (only M2/M4)

In simulation you can get this from PhysX contacts at the fingertips or gripper links.

Use a compact, learnable signal:

- Normal contact force magnitude per finger: $F_{\text{L}}, F_{\text{RF}}, F_{\text{L}}, F_{\text{R}}$ (2)
- Tangential magnitude per finger (optional): $T_{\text{L}}, T_{\text{RF}}, T_{\text{L}}, T_{\text{R}}$ (2)
- Binary contact indicators $c_{\text{L}}, c_{\text{RF}}, c_{\text{L}}, c_{\text{R}}$ (2)

- Optional: estimated slip proxy (relative tangential velocity at contact)

Recommended final force feature vector (8 dims):

$\text{oforce} = [\text{FL}, \text{FR}, \text{TL}, \text{TR}, \text{cL}, \text{cR}, \Delta F, \text{Fsum}]$

 $\text{o}_{\{\text{force}\}} = [\text{F_L}, \text{F_R}, \text{T_L}, \text{T_R}, \text{c_L}, \text{c_R}, \Delta F, \text{Fsum}]$
 $\text{F}_{\{\text{sum}\}} = [\text{FL}, \text{FR}, \text{TL}, \text{TR}, \text{cL}, \text{cR}, \Delta F, \text{Fsum}]$

Normalize forces by $\text{Fnorm}=20$ $\text{NF}_{\{\text{norm}\}}=20$ N (or your chosen limit).

Also add realistic noise:

- Force noise: $N(0, 0.2 \text{ N})$
 - Low-pass filter: $F_t = 0.8F_{t-1} + 0.2F_{\text{raw}}$
- $$F_t = 0.8F_{t-1} + 0.2F^{\{\text{raw}\}}$$

Actions (keep it stable)

Use an operational-space / Cartesian delta action (easier than direct torques):

Action vector:

- $\Delta x, \Delta y, \Delta z$ ($\Delta x, \Delta y, \Delta z$ (m per step, clipped e.g. $\pm 0.03 \text{ m}$)
- $\Delta r, \Delta p, \Delta y$ ($\Delta r, \Delta p, \Delta y$ (rad per step, clipped e.g. $\pm 0.15 \text{ rad}$)
- Gripper command $g \in [-1, 1]$ ($g \in [-1, 1]$ (close/open))

Then IsaacLab controller converts to joint targets/torques.

Policy architecture (publishable and robust)

You want something reviewers recognize as sensible, not fragile:

Encoder + fusion

- **Vision encoder:** small CNN (or ResNet-like) on 4 channels (RGBD)
 - Output 256-d embedding
- **Proprio/goal encoder:** 2-layer MLP \rightarrow 128-d
- **Force encoder:** 2-layer MLP \rightarrow 64-d (only for force models)
- Fuse by concatenation \rightarrow 256–384 dim
- Optional but strong: **GRU/LSTM (128)** after fusion
This is *very* relevant to “dropout duration” studies.

Actor-Critic

- PPO with Gaussian policy (continuous control)
 - **Asymmetric critic** (recommended): critic can see privileged state (object pose) to improve learning stability, while actor only sees sensors. This is widely accepted in sim RL.
-

Training algorithm and settings (practical defaults)

Use PPO (common in IsaacGym/IsaacLab workflows):

- $\gamma=0.99$ \gamma=0.99
- GAE $\lambda=0.95$ \lambda=0.95
- PPO clip $\epsilon=0.2$ \epsilon=0.2
- Learning rate: 3×10^{-4} \times 10^{-4}
- Entropy coef: 0.001 (or 0.0 if too stochastic)
- Value loss coef: 0.5
- Max grad norm: 1.0
- Rollout length: 32–64 steps
- Batch size: as large as GPU allows (vision limits env count)
- Envs: start with 32–128 (vision rendering is the bottleneck)

Train until plateau, but expect:

- M1 (clean vision) reaches >80% faster
 - robust variants may need 1.5–3× steps
-

Reward design (explicit, weighted, and consistent)

You asked for a specific reward/penalty system and weights. The key is: **stage-gated shaping** so the agent doesn't "cheat" by moving toward the goal before grasping.

Below is a reward that works well for pick-and-place and is easy to justify.

Definitions (computed from simulator state, not observations)

- $deo = \|pee-pobj\| d_{eo} = \|p_{ee}-p_{obj}\|$ deo = end-effector to object distance
- $dog = \|pobj-pgoal\| d_{og} = \|p_{obj}-p_{goal}\|$ dog = object to goal distance
- $h=zobjh = z_{obj} h=zobj$

- $hlift=0.10$ $h_{lift}=0.10$ $hlift=0.10$ m (lift threshold)
- aaa = action vector
- $Igrasp \in \{0,1\}$ $I_{grasp} \in \{0,1\}$ $Igrasp \in \{0,1\}$: object grasped (contact + gripper width condition)
- $Ilift \in \{0,1\}$ $I_{lift} \in \{0,1\}$ $Ilift \in \{0,1\}$: $h > h_{table} + hlith > h_{table} + h_{lift}$
- $Isucc \in \{0,1\}$ $I_{succ} \in \{0,1\}$ $Isucc \in \{0,1\}$: success condition
- $Idrop \in \{0,1\}$ $I_{drop} \in \{0,1\}$ $Idrop \in \{0,1\}$: object fell off / dropped (e.g., $z < z_{table} - 0.02$ $z < z_{table} - 0.02$ or out of bounds)
- $Icol \in \{0,1\}$ $I_{col} \in \{0,1\}$ $Icol \in \{0,1\}$: undesired collision (arm-table, self-collision)

Shaping terms

Use smooth exponentials:

- Reach:

$$rreach = \exp(-10 \cdot deo) r_{reach} = \exp(-10 \cdot d_{eo}) rreach = \exp(-10deo)$$

- Goal transport:

$$rgoal = \exp(-6 \cdot dog) r_{goal} = \exp(-6 \cdot d_{og}) rgoal = \exp(-6dog)$$

- Lift progress (linear, clipped):

$$rlift = \text{clip}(h - (h_{table} + 0.02)hlift, 0, 1) r_{lift} = \text{clip}(\frac{h - (h_{table} + 0.02)}{h_{lift}}, 0, 1) rlift = \text{clip}(hlith - (h_{table} + 0.02), 0, 1)$$

Full per-step reward (M1 / M3: vision+proprio)

```
rt=
1.0 rreach+Igrasp·1.5+Igrasp·2.0 rlift+Ilift·3.0 rgoal-0.01-0.002 ||a||2-1.0 Icol-5.0 Idrop\begin{aligned} r_t = & 1.0 \cdot r_{reach} + I_{grasp} \cdot 1.5 + I_{grasp} \cdot 2.0 \cdot r_{lift} \\ & + I_{lift} \cdot 3.0 \cdot r_{goal} - 0.01 - 0.002 \cdot \|a\|^2 - 1.0 \cdot I_{col} - 5.0 \cdot I_{drop} \end{aligned} rt=1.0rreach+Igrasp·1.5+Igrasp·2.0rlift+Ilift·3.0rgoal-0.01-0.002||a||2-1.0Icol-5.0Idrop
```

Terminal success bonus (and terminate episode):

$$rT+=10.0 \mathrel{\mathop{:}}= 10.0 \cdot I_{succ} rT+=10.0 I_{succ}$$

Force-aware reward addition (M2 / M4: +force)

You have two options:

Option A (recommended): Keep reward identical across models

This isolates the effect of adding force to observation. It's cleaner scientifically.

Option B (also publishable): add a small grasp-quality regularizer

This encourages using force without dominating learning.

Let $F_{sum} = FL + FR$, $F_{sum} = F_L + F_R$ (normal forces), and define a desired safe band:

- $F_{min} = 2F_{sum}$ ($F_{min} = 2$ N (too weak → slip likely))
- $F_{max} = 15F_{sum}$ ($F_{max} = 15$ N (too strong → “crushing” / unstable contact in sim))

Penalty only when grasped:

$$r_{force} = -0.05 \text{I}_{grasp}[\max(0, F_{min} - F_{sum}) + \max(0, F_{sum} - F_{max})] r_{force} = -0.05 \text{I}_{grasp}[\max(0, F_{min} - F_{sum}) + \max(0, F_{sum} - F_{max})]$$

Add:

$$rt += r_{force} \quad rt += r_{force}$$

This is small relative to the lift/goal rewards and mainly shapes contact behavior.

Publication advice: If you use Option B, you must also report Option A as an ablation, otherwise reviewers will say “it’s reward shaping, not sensing.”

Vision modality dropout: how to do it so the results are interpretable

You need two separate things:

1. **Dropout mechanism definition** (exactly how you corrupt observations)
2. **Evaluation protocol** (when dropout happens, how long, how often)

Dropout types (you should include at least 2)

Define “vision” = {RGB, Depth} jointly (main), and optionally separate them.

1. **Hard dropout (blackout):**
 - o RGB set to 0
 - o Depth set to d_{max} consistently
2. **Noise corruption:**

- RGB: add Gaussian noise + blur + brightness jitter
- Depth: Gaussian noise + missing pixels (set to $d_{max} \cdot \{max\}$)
Keep it bounded so it resembles sensor failure, not random TV static.

Important recommendation: include a “camera health flag” baseline

Real systems often know if a camera stream is missing. You can evaluate both:

- **Unannounced dropout:** no flag (harder)
- **Announced dropout:** add a 1-bit “vision_valid” to observations (more realistic)

This can be a strong discussion point.

Dropout schedule generator (concrete)

Do event-based dropout:

At each timestep:

- If not currently in dropout, start dropout with probability $p_{start} \cdot p_{start}$
- If started, sample duration LLL steps and apply corruption for LLL steps

Parameters:

- $p_{start}=0.02$ $p_{start}=0.02$ (roughly one dropout every ~50 steps on average)
- Durations $L \in \{0, 5, 10, 20, 40, 80\}$ $L \in \{0, 5, 10, 20, 40, 80\}$ steps
At 20 Hz: {0, 0.25s, 0.5s, 1s, 2s, 4s}

You can also sample LLL from a distribution (geometric or log-uniform) for smoother curves.

The “dropout duration affects success most” study (make it rigorous)

If you only vary duration, you’ll get messy results because the **timing** matters more than length.
A 0.5s dropout during grasp closure can be worse than a 2s dropout during transport.

Do this as a 2D sensitivity map

Evaluate success as a function of:

- **dropout onset time** (or phase)
- **dropout duration**

Two recommended variants:

Variant 1 (time-based grid)

- Onset bins: $t/T \in \{0-0.2, 0.2-0.4, 0.4-0.6, 0.6-0.8, 0.8-1.0\}$
- Durations: as above
- Report a heatmap of success rate

Variant 2 (phase-based, best for manipulation)

Define phases using privileged signals:

- Phase A: Reach (before first contact)
- Phase B: Grasp closure (first contact → grasp confirmed)
- Phase C: Lift/transport (grasp confirmed → near goal)
- Phase D: Place/release (near goal → released)

Then apply dropout starting at the phase boundary, varying duration.

This will almost certainly show:

- Vision dropout is most critical in Phase A/B
- Force helps most in Phase B/C/D
- Duration sensitivity differs by phase

That is a publishable result.

Quantify “most affected”

Don’t just eyeball curves—compute:

- $S(L)S(L)S(L)$: success vs duration
- “critical duration” $L_{50}L_{50}$: duration where success drops to 50% of clean-vision performance
- Fit a simple decay model per phase:

$$S(L) = S_\infty + (S_0 - S_\infty) \exp(-\alpha L)$$

Compare α between models/phases.

Training plan (what you should actually run)

Phase 1: Get M1 to >80% success (clean vision)

- Moderate domain randomization (don't overdo initially):
 - Object XY position
 - Object yaw
 - Mass/friction within small ranges
 - Camera jitter (small)
 - Lighting randomization (small)
- No dropout during training

Phase 2: Evaluate M1 under dropout (should drop)

Report:

- clean success (target >80%)
- dropout success vs L and vs phase
- noise vs hard-dropout

Phase 3: Train M2 (+force), same conditions as M1

- Same reward, same training steps budget, same seeds count
- Evaluate under dropout the same way

Phase 4 (recommended for publication): train robust baselines

Train M3 and optionally M4 with **vision corruption during training**:

- Start corruption only after N million steps (curriculum), or ramp probability:
 - p_{startp_start} ramp linearly $0 \rightarrow 0.02$
 - duration cap $0.25s \rightarrow 2s$ over training
- Keep corruption mild enough that learning still happens

This yields a compelling story:

- Clean-vision policies fail under intermittent sensor failures
- Robust-training helps
- Force sensing reduces brittleness further, especially in grasp/placement phases

Domain randomization (keep it realistic and targeted)

If your goal is robustness to **vision dropout**, you still need basic visual generalization:

Randomize:

- Textures (table/object)

- Lighting intensity/direction
- Camera extrinsics jitter (a few mm, $<1^\circ$)
- Depth noise model (Gaussian + missing pixels)
- Object friction/mass moderately

But do **not** randomize so aggressively that your clean-vision baseline can't learn—reviewers want controlled experiments.

Evaluation protocol (what to report so reviewers trust you)

Metrics (minimum)

- Success rate (%)
- Grasp success rate (%)
- Place success given grasp (%)
- Mean time-to-success (steps)
- Failure mode breakdown (drop, miss-grasp, collision, timeout)

Statistical practice (do this)

- ≥ 3 random seeds per method (5 is better)
- $\geq 500\text{--}2000$ evaluation episodes per condition (dropout durations multiply conditions; you can reduce by using confidence intervals smartly)
- Report 95% binomial confidence intervals on success rates

Compute fairness

Keep:

- same training steps
 - same network size except added force input encoder
 - same evaluation seeds/episodes
-

How to structure the paper (a complete writing guide)

Title ideas

- “Robust Vision-Based RL Manipulation Under Intermittent Sensor Failures via Gripper Force Sensing”
- “On the Sensitivity of Vision-Based Pick-and-Place Policies to Modality Dropout and the Role of Force Feedback”

Abstract (template)

1–2 sentences motivation (vision failures happen).

1 sentence method (train RGB-D+proprio RL; evaluate dropout; add force).

1 sentence results (force improves robustness, especially in grasp/place phases; quantify duration sensitivity).

1 sentence contribution (benchmark + protocol + findings).

1. Introduction

- Real robots lose frames / have glare / occlusions.
- Vision-only manipulation policies can be brittle.
- Force sensing is complementary: local contact info, grasp stability.
- Contributions list (bullet points).

2. Related Work

Organize into 3 buckets:

- Vision-based RL for manipulation
- Multi-modal sensing (vision + tactile/force)
- Robustness methods: modality dropout, domain randomization, sensor failure

3. Method

3.1 Environment (IsaacLab)

- Robot, object, camera placement, control frequency
- Observation definitions (RGB-D, proprio, force)
- Action space + controller

3.2 Policy architecture

- CNN for RGB-D, MLP for proprio/force, fusion, actor-critic
- Optional recurrence (if used)

3.3 Reward function

Include the exact equation and weights (use the ones above).

Explain stage-gating rationale.

3.4 Vision dropout modeling

- Hard dropout + noise dropout definitions
- Event-based dropout schedule
- Phase-based protocol

4. Experiments

- Training details (PPO hyperparameters, steps, env count, seeds)
- Baselines (M1–M4)
- Evaluation suite (durations, phases, announced/unannounced)

5. Results

Must include:

- Table: clean-vision success for all models
- Plot: success vs dropout duration (overall)
- Heatmap: success vs (phase or onset time) \times duration
- Ablation: force-only effect vs dropout-training effect
- Failure mode analysis (qualitative + quantitative)

6. Discussion

- Where force helps (grasp confirmation, slip mitigation, placement contact)
- Where it doesn't (object localization under full blackout pre-grasp)
- Practical implications: sensor redundancy, runtime detection of sensor failure

7. Limitations

Be honest:

- Sim contact models differ from real tactile sensors
- Rendering constraints limit env count
- Only single-object pick-and-place

8. Reproducibility

Include:

- Isaac Sim / IsaacLab version
- Full hyperparameters
- Random seeds
- Evaluation scripts for dropout schedules

Practical “don’t fail in week 3” implementation tips (IsaacLab + vision RL)

- **Vision rendering is the bottleneck.** Keep image small (84×84), reduce camera count.

- Use **frame skipping** cautiously (can destabilize grasp).
 - Normalize everything (proprio, forces, depth).
 - Start with **depth-only** debug runs if RGB training is unstable; then add RGB.
 - Add **action smoothing** (either in controller or via penalty) to avoid oscillations.
-

Expected outcomes (so you can sanity-check)

If you implement the above:

- **M1 (vision+proprio, clean training):**
 - Clean success: >80% achievable
 - Under dropout:
 - Huge drop if dropout hits Phase A/B
 - Moderate drop if dropout hits Phase C/D
- **M2 (+force, clean training):**
 - Clean success: similar to M1
 - Under dropout:
 - Improvement mainly in Phase B/C/D
 - Little/no improvement in Phase A
- **M3 (vision+proprio + dropout training):**
 - Clean success: slightly lower or similar
 - Under dropout: improved across many durations
- **M4 (+force + dropout training):**
 - Best robustness, especially long dropouts during grasp/placement

That pattern is coherent and defensible.

If you only do *one* thing to upgrade this to “publication-ready”

Do the **phase-based dropout** analysis (not just random duration). That is where your “duration affects most” question becomes a real result instead of a noisy curve.

If you want, I can also draft:

- a full **Methods** section text (ready-to-paste),
- an **Experiment table layout** (conditions × durations × phases),
- and pseudocode for the dropout scheduler and phase detector (in a paper-friendly format).

